

# Introduction to Statistics and Data Visualisation with R

## Lausanne, January 2026

Joao Lourenço and Rachel Marcone

## Clustering, Heatmap and PCA



# Visual representation

# Dimension

Dimension:  
the number of coordinates we need to  
locate a point in a given space.

# 2-dimension

	Coordinate1	Coordinate2
x1	1	2
x2	1	4
...	...	...
xN	5	2

1	2
..	..
..	..


*Two dimensions: latitude and longitude*



# 3-dimension

	Coordinate1	Coordinate2	Coordinate3	1	2	3
x1	1	2	5			
x2	1	4	7			
...	...	...	...	...	...	...
xN	5	2	1			

*Three dimensions: latitude, longitude and altitude*




# n-dimension

	Coordinate1	Coordinate2	Coordinate3	...	Coordinate-n	1	2	3	n
x1	1	2	5	...	18				
x2	1	4	7	...	2				
...	...	...	...	...	...	..	..	..	..
xN	5	2	1	...	4				

Dimension in biology?


# Genes

# 3-dimension




# n-dimension

?




# Dimension reduction


## Starting point: Big Data





**Starting point: Big Data**



**End result: human readable**








clarity of  
representation

Over-simplification




**There are many possibilities and  
there is not a « better » one than  
another. It depends on what you want  
to show.**

# World Map



<https://ontheworldmap.com/>





<https://www.shutterstock.com/fr/image-vector/world-map-pacific-china-asia-centered-1731018682>

# Principal Component Analysis (PCA)

Pearson (1901) and Hotelling (1933)

- PCA is based on variance
- PCA is the best angle to see and evaluate the data


Which and how?




# PCA- Principal component analysis

1. Largest variance first

# PCA- Principal component analysis




# PCA- Principal component analysis




# PCA- Principal Component Analysis

2. Select uncorrelated principal axis  
(orthogonal)


# PCA- Principal Component Analysis




# PCA- Principal Component Analysis




# PCA- Principal Component Analysis




# PCA- Principal Component Analysis








2. Select uncorrelated principal axis  
(orthogonal)





centroid












**Without centroid**



# The PCA axis

- The PC are linear combination of the original axis.
- The estimated parameters of the linear combination is known and therefore we can know positively or negatively how much it goes into one direction or the other one.
- Indeed as the original axis are  $g_1, g_2, g_3 \dots$  and the new axis are  $a_1g_1 + a_2g_2 \dots$ , one takes the  $a_i$  that are the highest, positively and negatively and therefore knows which features are mostly representing the axis you see.
- Observation : **Scaling** is important, if one variable is on a different scale than another, it will dominate the PCA procedure as the largest variance might be observed there, and the low dimension plot will really just be visualizing that dimension.


# Mathematically

You calculate the covariance matrix meaning a matrix containing two-by-two covariances

- if positive then : the two variables increase or decrease together (correlated)
- if negative then : One increases when the other decreases (Inversely correlated)
- $(\text{Corr}(X,Y) = \text{Cov}(X,Y)/\text{sd}(x)*\text{sd}(y))$
- Eigenvectors and eigenvalues are the linear algebra concepts that we need to compute from the covariance matrix in order to determine the ***principal components*** of the data.

# Mathematically

- eigenvectors of the Covariance matrix are *the directions of the axes where there is the most variance (this is something you can prove mathematically!)*
- eigenvalues are the coefficients attached to eigenvectors, which give the *amount of variance carried in each Principal Component.*
- After having the principal components, to compute the percentage of variance (information) accounted for by each component, we divide the eigenvalue of each component by the sum of eigenvalues.



Scree Plot for Genetic Data. ([Source](#).)

<https://towardsdatascience.com/a-one-stop-shop-for-principal-component-analysis-5582fb7e0a9c>

# How many PCs?

- Method 1: We **arbitrarily select** a number of principal components to include. Suppose I wanted to keep five principal components in my model. In the genetic data case above, these five principal components explains about 66% of the total variability that would be explained by including all 13 principal components.
- Method 2: Suppose I wanted to include **enough principal components to explain 90%** of the total variability explained by all 13 principal components. In the genetic data case above, I would include the first 10 principal components and drop the final three variables from  $Z^*$ .
- Method 3: Here, we want to “**find the elbow**.” In the scree plot above, we see there’s a big drop in proportion of variability explained between principal component 3 and the following. In this case, we’d likely include the first three features and drop the remaining features. As you can see, this method is a bit subjective as “elbow” doesn’t have a mathematically precise definition and, in this case, we’d include a model that explains only about 42% of the total variability.

# In R

```
>pca<-prcomp(data, center = TRUE, scale. = FALSE)  
#coordinate of sample on components were identified
```

## **#Importance of components**

```
>summary(pca)
```

# In R

```
>pca<-prcomp(data, center = TRUE, scale. = FALSE)
#coordinate of sample on components were identified
```

## #Importance of components

```
>summary(pca)
```

```
>pca$x
```

```
>plot(pca$x)
```

	PC1	PC2	PC3	PC4
[1,]	-2.684125626	-0.319397247	0.027914828	0.0022624371
[2,]	-2.714141687	0.177001225	0.210464272	0.0990265503
[3,]	-2.888990569	0.144949426	-0.017900256	0.0199683897
[4,]	-2.745342856	0.318298979	-0.031559374	-0.0755758166
[5,]	-2.728716537	-0.326754513	-0.090079241	-0.0612585926
[6,]	-2.280859633	-0.741330449	-0.168677658	-0.0242008576
[7,]	-2.820537751	0.089461385	-0.257892158	-0.0481431065
[8,]	-2.626144973	-0.163384960	0.021879318	-0.0452978706
[9,]	-2.886382732	0.578311754	-0.020759570	-0.0267447358
[10,]	-2.672755798	0.113774246	0.197632725	-0.0562954013
[11,]	-2.506947091	-0.645068899	0.075318009	-0.0150199245
[12,]	2.612255231	0.014720030	0.102150260	-0.1563702070


# Center, Scale

- Why should we center, why should we scale ?
- Will see this through the exercises
- In principal if one considers PCA to be the Eigenvalue of the covariance matrix, then centering yes or no will not change the result.
- In prcomp, however, "PCA" is defined as computing the eigenvalues of the  $X^T X / (n-1)$  matrix, which in a centered data is exactly the covariance matrix, otherwise not.

# Center, Scale


- Although proven not to be exactly true \*, this will result in first PCs capturing the mean of the data as this "explains" most of the variance in the model.
- The scaling will determine if you compute eigenvectors on the covariance matrix (if unscaled) or on the correlation matrix (if scaled).
- This again (mostly) means that what you will capture in the first PCs is mostly what is in a bigger scale.
- \*The Effect of Data Centering on PCA Models Neal B. Gallagher, Donal O'Sullivan, Manuel Palacios

# The perfect human is Puerto Rican...



<https://liorpachter.wordpress.com/2014/12/02/the-perfect-human-is-puerto-rican/>

# ... or an alien ?



# What's new ?


PCA	linear	Matrix Factorization		
ICA	linear	Matrix Factorization		
MDS	non-linear	Matrix Factorization		
Sparce NMF	non-linear	Matrix Factorization	2010	<a href="https://pdfs.semanticscholar.org/664d/40258f12ad28ed0b7dc272935ad72a150db.pdf">https://pdfs.semanticscholar.org/664d/40258f12ad28ed0b7dc272935ad72a150db.pdf</a>
cPCA	non-linear	Matrix Factorization	2018	<a href="https://doi.org/10.1038/s41467-018-04608-8">https://doi.org/10.1038/s41467-018-04608-8</a>
ZIFA	non-linear	Matrix Factorization	2015	<a href="https://doi.org/10.1186/s13059-015-0805-z">https://doi.org/10.1186/s13059-015-0805-z</a>
ZINB-WaVE	non-linear	Matrix Factorization	2018	<a href="https://doi.org/10.1038/s41467-017-02554-5">https://doi.org/10.1038/s41467-017-02554-5</a>

Diffusion maps	non-linear	graph-based	2005	<a href="https://doi.org/10.1073/pnas.0500334102">https://doi.org/10.1073/pnas.0500334102</a>
Isomap	non-linear	graph-based	2000	<a href="https://doi.org/10.1126/science.290.5500.2319">10.1126/science.290.5500.2319</a>
t-SNE	non-linear	graph-based	2008	<a href="https://lvdmaaten.github.io/publications/papers/JMLR_2008.pdf">https://lvdmaaten.github.io/publications/papers/JMLR_2008.pdf</a>
- BH t-SNE	non-linear	graph-based	2014	<a href="https://lvdmaaten.github.io/publications/papers/JMLR_2014.pdf">https://lvdmaaten.github.io/publications/papers/JMLR_2014.pdf</a>
- Flt-SNE	non-linear	graph-based	2017	<a href="https://arxiv.org/abs/1712.09005">arXiv:1712.09005</a>
LargeVis	non-linear	graph-based	2018	<a href="https://arxiv.org/abs/1602.00370">arXiv:1602.00370</a>
UMAP	non-linear	graph-based	2018	<a href="https://arxiv.org/abs/1802.03426">arXiv:1802.03426</a>
PHATE	non-linear	graph-based	2017	<a href="https://www.biorxiv.org/content/biorxiv/early/2018/06/28/120378.full.pdf">https://www.biorxiv.org/content/biorxiv/early/2018/06/28/120378.full.pdf</a>

scvis	non-linear	Autoencoder (MF)	2018	<a href="https://doi.org/10.1038/s41467-018-04368-5">https://doi.org/10.1038/s41467-018-04368-5</a>
VASC	non-linear	Autoencoder (MF)	2018	<a href="https://doi.org/10.1016/j.gpb.2018.08.003">https://doi.org/10.1016/j.gpb.2018.08.003</a>


# Clustering

# Clustering



**Point cloud**


# Clustering





# Clustering

Clustering method are divided into two categories :

Partitioning clustering



Hierarchical clustering





\*Handbook of cluster analysis, Hennig C. et al.

# Partitioning clustering


**Convex partitioning. Example: K-means**

**Density based approaches. Example: DBSCAN**


**Model-based approaches. Example: Mclust**




# Hierarchical Clustering




# Hierarchical Clustering



# Hierarchical Clustering



# Hierarchical Clustering



# Distance

# Euclidean

$$X = (2, 0)$$

$$Y = (-2, -2)$$


$$\sqrt{\sum (y - x)^2}$$

$$= \sqrt{(-2 - 2)^2 + (-2 - 0)^2}$$

$$= \sqrt{4^2 + 2^2}$$

$$= \sqrt{20}$$

$$= 4.47$$



It represents the “multivariate dissimilarity” of X & Y

# Squared Euclidean

$$X = (2, 0)$$


$$Y = (-2, -2)$$

$$\Sigma (y - x)^2$$

$$= ((-2 - 2)^2 + (-2 - 0)^2 )$$

$$= (4^2 + 2^2)$$

$$= 20$$



It represents the “multivariate dissimilarity” of X & Y

# City Block

$$X = (2, 0)$$


$$Y = (-2, -2)$$

$$\Sigma |y - x|$$

$$= (|-2 - 2| + |-2 - 0|)$$

$$= |-4| + |-2|$$

$$= 6$$



# Distance Measures in 2D

- Euclidean  $\sqrt{[\sum (y - x)^2]}$
- Squared Euclidean  $\sum (y - x)^2$
- City-Block  $\sum |y - x|$

# Distance Measures in nD

- Euclidean

$$d = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots + (a_n - b_n)^2}$$

- Squared Euclidean

$$d = (a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots + (a_n - b_n)^2$$


- City-Block

$$d = |a_1 - b_1| + |a_2 - b_2| + \dots + |a_n - b_n|$$


# In R

>?dist


# Distance matrix



# Distance matrix



Heatmap of distance matrix



Sample 1

Sample 2

Sample 3

Sample 4

Sample 5

Sample 6

Sample 1

Sample 2

Sample 3

Sample 4



Sample 5

Sample 6

# In R


```
>?heatmap  
>heatmap(distanceMatrix, Colv=NA, Rowv=NA,  
scale="none")
```

# How to aggregate clusters?




Which clusters to combine?

# Single linkage




Distance between closest elements in clusters

# Complete Linkage




Distance between furthest elements in clusters

# Average Linkage




Average of all pairwise distances

# Centroid Condensation (mean)



Distance between centroids (means) of two clusters


# Median Condensation




Distance between median distances of two clusters

# Clustering methods


# Hierarchical Clustering



At the beginning every point is a cluster in it self, then we agglomerate ...




## Euclidean distance




Euclidean distance  
complete Linkage

# **Determine the Termination Condition (TC)**



Euclidean distance  
complete Linkage

Euclidean distance  
complete Linkage



# Let's practice In R

**#create a random matrix**

```
>mat <- matrix(data = rnorm(300, mean= 100,  
sd=10), nrow = 150, ncol = 2)
```

**#Euclidian distance**

```
>mat.dist<-as.matrix(dist(mat))
```

**#show heatmap**

```
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none")
```

**#change heatmap's color**

```
>colorScale <- colorRampPalette(c("blue",  
"green", "yellow", "red", "darkred"))(1000)  
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none",  
col=colorScale)
```

# Let's practice In R

```
#create a random matrix
```

```
>mat <- matrix(data = rnorm(300, mean= 100,  
sd=10), nrow = 150, ncol = 2)
```

```
#Euclidian distance
```

```
>mat.dist<-as.matrix(dist(mat))
```

```
#show heatmap
```

```
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none")
```

```
#change heatmap's color
```

```
>colorScale <- colorRampPalette(c("blue",
```

```
"green", "yellow", "red", "darkred"))(1000)
```

```
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none",  
col=colorScale)
```

# Let's practice In R

```
#create a random matrix
```

```
>mat <- matrix(data = rnorm(300, mean= 100,  
sd=10), nrow = 150, ncol = 2)
```

```
#Euclidian distance
```

```
>mat.dist<-as.matrix(dist(mat))
```

```
#show heatmap
```

```
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none")
```

```
#change heatmap's color
```

```
>colorScale <- colorRampPalette(c("blue",  
"green","yellow","red","darkred"))(1000)
```

```
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none",  
col=colorScale)
```

# Let's practice In R

**#create a random matrix**

```
>mat <- matrix(data = rnorm(300, mean= 100,  
sd=10), nrow = 150, ncol = 2)
```

**#Euclidian distance**

```
>mat.dist<-as.matrix(dist(mat))
```

**#show heatmap**

```
>heatmap(mat.dist,Colv=NA, Rowv=NA, scale="none")
```

**#change heatmap's color**

```
>colorScale <- colorRampPalette(c("blue",  
"green","yellow","red","darkred"))(1000)  
>heatmap(mat.dist,Colv=NA, Rowv=NA, scale="none",  
col=colorScale)
```

# Let's practice In R

```
#create a random matrix
```

```
>mat <- matrix(data = rnorm(300, mean= 100,  
sd=10), nrow = 150, ncol = 2)
```

```
#Euclidian distance
```

```
>mat.dist<-as.matrix(dist(mat))
```

```
#show heatmap
```

```
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none")
```

```
#change heatmap's color
```

```
>colorScale <- colorRampPalette(c("blue",  
"green","yellow","red","darkred"))(1000)
```

```
>heatmap(mat.dist, Colv=NA, Rowv=NA, scale="none",  
col=colorScale)
```

# How to do hierarchical clustering in R?

```
>?hclust
```

```
#Euclidian distance
```

```
>distE<-dist(mat)
>mat.distE<-as.matrix(dist(mat))
>heatmap(mat.distE, Colv=NA, Rowv=NA, scale="none")
>hE<-hclust(distE,"complete")
>plot(hE)
```

```
#manhattan distance
```

```
>distC<-dist(mat,method="manhattan"
>mat.distC<-as.matrix(dist(mat,method="manhattan"))
>heatmap(mat.distC, Colv=NA, Rowv=NA, scale="none")
>hC<-hclust(distC,"complete")
>plot(hC)
```

# How to do hierarchical clustering in R?

```
>?hclust
```

## #Euclidian distance

```
>distE<-dist(mat)
>mat.distE<-as.matrix(dist(mat))
>heatmap(mat.distE, Colv=NA, Rowv=NA, scale="none")
>hE<-hclust(distE,"complete")
>plot(hE)
```

## #manhattan distance

```
>distC<-dist(mat,method="manhattan")
>mat.distC<-as.matrix(dist(mat,method="manhattan"))
>heatmap(mat.distC, Colv=NA, Rowv=NA, scale="none")
>hC<-hclust(distC,"complete")
>plot(hC)
```

# How to do hierarchical clustering in R?

```
>?hclust
```

## #Euclidian distance

```
>distE<-dist(mat)
>mat.distE<-as.matrix(dist(mat))
>heatmap(mat.distE, Colv=NA, Rowv=NA, scale="none")
>hE<-hclust(distE,"complete")
>plot(hE)
```

## #manhattan distance

```
>distC<-dist(mat,method="manhattan")
>mat.distC<-as.matrix(dist(mat,method="manhattan"))
>heatmap(mat.distC, Colv=NA, Rowv=NA, scale="none")
>hC<-hclust(distC,"complete")
>plot(hC)
```

# How to do hierarchical clustering in R?

```
>?hclust
```

## #Euclidian distance

```
>distE<-dist(mat)
>mat.distE<-as.matrix(dist(mat))
>heatmap(mat.distE, Colv=NA, Rowv=NA, scale="none")
>hE<-hclust(distE,"complete")
>plot(hE)
```


## #manhattan distance

```
>distC<-dist(mat,method="manhattan")
>mat.distC<-as.matrix(dist(mat,method="manhattan"))
>heatmap(mat.distC, Colv=NA, Rowv=NA, scale="none")
>hC<-hclust(distC,"complete")
>plot(hC)
```

# K-means Clustering


Number of clusters = 3

a




Start with  
3 initial  
points

b




For each  
point  
determine to  
which initial  
point it is the  
closest

c



Move initial  
points to the  
centroids of  
the clusters

d




Color again  
each point!  
Repeat b and  
c until  
obtaining  
stabilisation


# K-means Clustering

Number of clusters = 3


a





b



d



c



# K-means & C-means

Drawbacks:

1. Specify number of clusters
2. Non probabilistic methods
3. Not stable

# Kmeans in R

```
>mat <- matrix(data = rnorm(300, mean=100, sd=10),  
                 nrow = 150,  
                 ncol = 2)  
>df<-data.frame(x)  
  
>kmeans(df, 3)
```


# Kmeans in R


```
>mat <- matrix(data = rnorm(300, mean= 100, sd=10),  
                 nrow = 150,  
                 ncol = 2)  
>df<-data.frame(mat)  
  
>kmeans(df,3)  
  
>cl.1 <- kmeans(df, 3, iter.max = 1)  
>plot(df, col = cl.1$cluster)  
>points(cl.1$centers, col = 1:5, pch =  
8)
```


# Kmeans in R

```
>mat <- matrix(data = rnorm(300, mean= 100, sd=10),  
                 nrow = 150,  
                 ncol = 2)  
>df<-data.frame(mat)  
  
>kmeans(df,3)  
  
>cl.1 <- kmeans(df, 3, iter.max = 1)  
>plot(df, col = cl.1$cluster)  
>points(cl.1$centers, col = 1:5, pch = 8)  
  
>cl.10 <- kmeans(df, 3, iter.max = 10)  
>plot(df, col = cl.10$cluster)  
>points(cl.10$centers, col = 1:5, pch = 8)  
  
>cl.100 <- kmeans(df, 3, iter.max = 100)  
>plot(df, col = cl.100$cluster)  
>points(cl.100$centers, col = 1:5, pch = 8)
```


# Model-based Clustering








distribution  
(univariate, spherical, diagonal,  
ellipsoidal)

















data volume  
(equal, variable)



shape  
(equal, variable)

# Volume, Shape, Orientation



# Model selection

identifier	Model	HC	EM	Distribution	Volume	Shape	Orientation
E		•	•	(univariate)	equal		
V		•	•	(univariate)	variable		
EII	$\lambda I$	•	•	Spherical	equal	equal	NA
VII	$\lambda_k I$	•	•	Spherical	variable	equal	NA
EEI	$\lambda A$		•	Diagonal	equal	equal	coordinate axes
VEI	$\lambda_k A$		•	Diagonal	variable	equal	coordinate axes
EVI	$\lambda A_k$		•	Diagonal	equal	variable	coordinate axes
VVI	$\lambda_k A_k$		•	Diagonal	variable	variable	coordinate axes
EEE	$\lambda DAD^T$	•	•	Ellipsoidal	equal	equal	equal
EEV	$\lambda D_k AD_k^T$		•	Ellipsoidal	equal	equal	variable
VEV	$\lambda_k D_k AD_k^T$		•	Ellipsoidal	variable	equal	variable
VVV	$\lambda_k D_k A_k D_k^T$	•	•	Ellipsoidal	variable	variable	variable

BIC= Bayesian information criterion

# BIC

- the model with **the highest BIC** is preferred
- BIC is a function of the number of parameters of the model
- The goodness of the fit of the model
- The sample size

Number of  
parameters

Best likelihood




# Mclust In R

```
>?mclustBIC  
>?Mclust  
  
>BIC <- mclustBIC(df)  
>plot(BIC)  
>summary(BIC)  
>mod1 <- Mclust(df, x = BIC)  
>summary(mod1, parameters = TRUE)  
>plot(mod1, what = "classification")
```

**Once you have the clusters,  
what do you do with them ?**

**PC2 (12.8%)**

**PC1 (20.5%)**



# Challenge

## Points in plates

1. Import the data from dataClustering.csv
2. What is the dimension of this dataset?
3. How many data point do we have?
4. Evaluate Euclidean distance of points in a plates
5. Classify points to find clusters using hierarchical clustering and the average agglomeration method

# Challenge

## Points in plates-continuous

6. We expect to have 3 clusters. When you apply k-means algorithm using 1 iteration, does it differ from applying it using 10 or 100 iterations?

7. What is the outcome of the C-means clustering?

```
install.packages("e1071")
```

```
library(e1071)
```

```
?cmeans
```

# Challenge

Points in plates-continuous

Library(mclust)

8. What are the top 3 models *mclustBIC* function suggests based on the BIC criterion?
9. How many clusters did it find using the top model?
10. Plot the outcome

Thank you for your attention